

Linear and Softmax Layers in BERT Fine-tuning: Formulas and Functions

In BERT fine-tuning, the **linear and softmax layers** form the classification head that processes the final hidden representations from the transformer to produce predictions. Here's a comprehensive explanation of their mathematical formulas and functions:

Linear Layer (Classification Head)

The linear layer is a fully connected layer that transforms the BERT output representation into class scores.

Formula:

$$z = W \cdot h + b$$

Where:

- z is the output logits (raw prediction scores)
- W is the weight matrix with dimensions $\text{num_classes} \times \text{hidden_size}$
- h is the hidden representation from BERT (typically the [CLS] token representation)
- b is the bias vector with dimensions num_classes

What it does:

- **Transforms representations:** Converts the high-dimensional BERT output (usually 768 dimensions for BERT-base) into class-specific scores^{[1] [2]}
- **Learns task-specific mappings:** The weight matrix W learns which features in the BERT representation are most important for each class
- **Provides linear transformation:** Each output neuron computes a weighted sum of all input features plus a bias term

Softmax Layer (Activation Function)

The softmax function converts the raw logits from the linear layer into probability distributions.

Formula:

$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

Where:

- z_i is the logit for class i
- K is the total number of classes
- e is the exponential function

What it does:

- **Normalizes outputs:** Ensures all class probabilities sum to 1^[3] ^[4]
- **Amplifies differences:** Uses exponential function to emphasize higher-scoring classes
- **Enables probabilistic interpretation:** Converts raw scores into meaningful probabilities for decision-making

Combined Process in BERT Fine-tuning

The complete classification pipeline follows this sequence:

1. **BERT Processing:** Input text → BERT layers → [CLS] token representation h
2. **Linear Transformation:** $z = W \cdot h + b$
3. **Softmax Activation:** $p = \text{softmax}(z)$
4. **Prediction:** $\hat{y} = \arg\max(p)$

Training Process

During fine-tuning, these layers are optimized using:

- **Cross-entropy loss:** $L = -\sum_{i=1}^K y_i \log(p_i)$
- **Backpropagation:** Gradients flow back through softmax → linear → BERT layers
- **Parameter updates:** Both the linear layer weights and BERT parameters are updated

Practical Implementation

In practice, the linear and softmax layers work together as follows^[2] ^[5]:

```
# Conceptual implementation
linear_output = torch.matmul(bert_output, weight_matrix) + bias
probabilities = torch.softmax(linear_output, dim=-1)
```

Key Functions

Linear Layer Functions:

- Feature extraction and dimensionality reduction
- Task-specific weight learning
- Bias adjustment for class imbalances

Softmax Layer Functions:

- Probability normalization
- Confidence scoring
- Multi-class decision making

Variations and Enhancements

Recent research has explored alternatives to the standard linear+softmax approach:

- **Multiple linear layers:** Adding depth to the classification head^[2]
- **Attention mechanisms:** Using attention to weight different parts of the representation^[4]
- **Alternative activation functions:** Exploring functions beyond softmax for specific tasks^[6]

The linear and softmax layers, while conceptually simple, are crucial for translating BERT's rich contextual representations into task-specific predictions. Their mathematical simplicity belies their importance in the overall fine-tuning process, serving as the bridge between pre-trained language understanding and downstream task performance.



1. <https://www.semanticscholar.org/paper/b964afe5b755022f1f1e6915d23df9a7f65c911c>
2. <https://ieeexplore.ieee.org/document/10742347/>
3. <https://www.mdpi.com/1424-8220/23/3/1481>
4. <https://journals.sagepub.com/doi/10.3233/KES-230066>
5. <https://www.mdpi.com/2078-2489/14/8/467>
6. <https://arxiv.org/abs/2408.08803>